

Exercise Set 1

Problem 1

Task a

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df = pd.read_csv("train.csv")
```

```
df.drop(columns=['id', 'partlybad'], inplace=True)
```

Task b

```
df[['T84.mean', 'UV_A.mean', 'CS.mean']].describe()
```

	T84.mean	UV_A.mean	CS.mean
count	450.000000	450.000000	450.000000
mean	6.439594	10.797330	0.003083
std	9.827282	6.626678	0.002246
min	-23.627710	0.438965	0.000343
25%	-0.733160	4.304971	0.001436
50%	7.661622	11.506937	0.002548
75%	14.118346	16.573188	0.004119
max	27.001804	22.500037	0.013706

Task c

```
arr = df['T84.mean'].values
mean = arr.mean()
std = arr.std()
print(f"Mean: {mean}, Std: {std}")
```

Mean: 6.439594405677285, Std: 9.816356767300478


Task d

```
fig, axes = plt.subplots(1, 2, figsize=(8, 3))

axes[0].bar(df['class4'].value_counts().index,
df['class4'].value_counts().values)
axes[0].set_title('Class 4')
axes[0].set_xlabel('Class 4')
axes[0].set_ylabel('Count')


axes[1].hist(df['CO242.mean'])
axes[1].set_title('CO242 Mean')
axes[1].set_xlabel('CO242 Mean')
axes[1].set_ylabel('Count')

plt.tight_layout()
plt.show()
```



Task e

```
sns.pairplot(df[['UV_A.mean', 'T84.mean', 'H2084.mean']])
plt.show()
```



Task f

```
most_common_class = df['class4'].value_counts().idxmax()
p = (len(df) - df['class4'].value_counts()['nonevent']) / len(df)

test = pd.read_csv("test.csv")
submission = pd.DataFrame({'id': test['id'], 'class4': most_common_class, 'p':
p})
submission.to_csv("submission.csv", index=False)
```

Problem 2

Task a

```
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import PolynomialFeatures, StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import KFold, cross_val_score
```

```
train_syn = pd.read_csv("train_syn.csv")
valid_syn = pd.read_csv("valid_syn.csv")
test_syn = pd.read_csv("test_syn.csv")
```

```
degrees = np.arange(0, 9)
models = {}
```

```
train_loss = []
valid_loss = []
test_loss = []
test_trva_loss = []
cv_loss = []
```

```
X_train = train_syn.drop(columns=['y'])
y_train = train_syn['y']
X_valid = valid_syn.drop(columns=['y'])
y_valid = valid_syn['y']
X_test = test_syn.drop(columns=['y'])
y_test = test_syn['y']
X_trva = pd.concat([X_train, X_valid], axis=0)
y_trva = pd.concat([y_train, y_valid], axis=0)
```

```
def make_poly_ols(degree: int) -> Pipeline:
    include_bias = True if degree == 0 else False
    return Pipeline([
        ("poly", PolynomialFeatures(degree=degree,
    include_bias=include_bias)),
        ("scaler", StandardScaler()),
        ("ols", LinearRegression())
    ])
```

```
def compute_loss(degree: int, X_train: pd.DataFrame, y_train: pd.Series,
X_test: pd.DataFrame, y_test: pd.Series) -> float:
    model = make_poly_ols(degree)
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    return np.mean((y_test - y_pred) ** 2)
```

```

def cv_score_model(model: Pipeline, X: pd.DataFrame, y: pd.Series, cv: int =
10) -> float:
    kf = KFold(n_splits=cv, shuffle=True, random_state=42)
    scores = cross_val_score(model, X, y, scoring='neg_mean_squared_error',
cv=kf)
    return -scores.mean()

def compute_all_losses(degree: int):
    train_loss.append(compute_loss(degree, X_train, y_train, X_train,
y_train))
    valid_loss.append(compute_loss(degree, X_train, y_train, X_valid,
y_valid))
    test_loss.append(compute_loss(degree, X_train, y_train, X_test, y_test))
    test_trva_loss.append(compute_loss(degree, X_trva, y_trva, X_test,
y_test))

    cv_loss.append(cv_score_model(make_poly_ols(degree), X_trva, y_trva))
    model = make_poly_ols(degree)
    model.fit(X_trva, y_trva)
    models[degree] = model

```

```

for degree in degrees:
    compute_all_losses(degree)

results_df = pd.DataFrame({
    'Degree': degrees,
    'Train': train_loss,
    'Valid': valid_loss,
    'Test': test_loss,
    'TRVA': test_trva_loss,
    'CV': cv_loss
}).set_index('Degree')
print(results_df)

```

Degree	Train	Valid	Test	TRVA	CV
0	18.459585	32.342363	22.100658	21.620221	28.779634
1	4.088535	7.127844	8.876307	9.934943	7.252893
2	0.218586	0.293732	0.245847	0.214016	0.321785
3	0.216819	0.283447	0.290080	0.275111	0.386635
4	0.118796	0.624726	0.969078	0.223993	0.641643
5	0.096532	0.573480	4.894836	1.039089	0.393977
6	0.007574	3.416787	213.297140	0.881471	0.399130
7	0.004999	6.862993	1261.988070	0.271719	1.738198
8	0.002082	401.651780	154266.871360	11.223554	25.777011

Polynomial with the lowest CV loss is the polynomial which probably generalizes the best. So that's the one to choose.

Task b


```
p = [0, 1, 2, 3, 4, 8]
x = np.linspace(start=-3, stop=3, num=256)
x = x.reshape(-1, 1)
x = pd.DataFrame(x, columns=['x'])
```

```
fig, axes = plt.subplots(2, 3, figsize=(15, 8))

for i, deg in enumerate(p):
    row = i // 3
    col = i % 3
    ax = axes[row, col]

    model = models[deg]
    y_pred = model.predict(x)

    ax.scatter(X_trva, y_trva, color='blue', label='Data', alpha=0.5)
    ax.plot(x, y_pred, color='red', label=f'Polynomial Degree {deg}')
    ax.set_title(f'Polynomial Regression (Degree {deg})')
    ax.set_xlabel('x')
    ax.set_ylabel('y')
    ax.legend()
```



Task c

```
from sklearn.dummy import DummyRegressor
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error as mse
from sklearn.linear_model import Ridge

def rmse(y_true, y_pred):
    return np.sqrt(mse(y_true, y_pred))
```

```
train = pd.read_csv("train_real.csv")
test = pd.read_csv("test_real.csv")

X_train = train.drop(columns=['Next_Tmax'])
y_train = train['Next_Tmax']
X_test = test.drop(columns=['Next_Tmax'])
y_test = test['Next_Tmax']

train_loss = []
test_loss = []
cv_loss = []

models = {
    'Dummy': DummyRegressor(),
    'OSL Linear Regression': LinearRegression(),
    'Random Forest': RandomForestRegressor(),
    'SVR': Pipeline([('scaler', StandardScaler()), ('svr', SVR())]),
    'Ridge': Ridge()
}
```

```
def get_losses(model):
    model.fit(X_train, y_train)
    y_pred_test = model.predict(X_test)
    y_pred_train = model.predict(X_train)
    train_loss.append(rmse(y_train, y_pred_train))
    test_loss.append(rmse(y_test, y_pred_test))

    kf = KFold(n_splits=10, shuffle=True, random_state=42)
    scores = cross_val_score(model, X_train, y_train,
                              scoring='neg_root_mean_squared_error', cv=kf)
    cv_loss.append(-scores.mean())
```

```
for name, model in models.items():
    get_losses(model)
```

```

results_df = pd.DataFrame({
    'Model': list(models.keys()),
    'Train RMSE': train_loss,
    'Test RMSE': test_loss,
    'CV RMSE': cv_loss
}).set_index('Model')

print(results_df)

```

	Train RMSE	Test RMSE	CV RMSE
Model			
Dummy	3.089275	2.997846	3.080440
OSL Linear Regression	1.378378	1.483817	1.447004
Random Forest	0.528063	1.469373	1.440465
SVR	1.406482	1.644513	1.727787
Ridge	1.379322	1.484116	1.443108

1. The lowest test_RMSE is achieved by the Linear Regression model. So this is the best one for the given test set. However Random Forest has the best CV_RMSE, so it might generalize better to unseen data.
2. We can see that Generally Train RMSE is lower than Test RMSE, which is expected as models tend to perform better on the data they were trained on. CV RMSE is usually somewhere in between, as it estimates the performance better due to cross-validation.
3. We can achieve better performance by hyperparameter tuning, feature engineering, or trying more advanced models.

Problem 3

Task a

	Small Models	Flexibility	Average Models	Flexibility	High Flexibility Models
Training Error	Higher - simple to fit data	model too	Lower - captures main structure		Lowest - can fit almost perfectly
Test Error	Higher - underfitting		Lowest - best balance		Higher - overfitting
(squared) Bias	High - oversimplified assumptions		Medium		Low - fits data closely
Variance	Low - stable predictions		Medium		High - sensitive to data noise

	Small Models	Flexibility	Average Models	Flexibility	High Flexibility Models
irreducible Error	Constant - noise in data cannot be reduced by any model		Constant		Constant

Task b